

Code

Issues2


Pull requests5

Actions

Projects



Security


Insights


 master ▾



deep-learning-keras-tf-tutorial / 10_gpu_benchmarking /
gpu_performance_test_small_image_classification.ipynb

 **Karandeep Grover** ENH: moved all the active tutorial content from py/deeplearning 

 0 contributors

735 lines (735 sloc) | 54.4 KB 

Small Image Classification Using Simple Artificial Neural Network: GPU Benchmarking

```
In [4]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np
# Version Information
# tensorflow 2.2.0 , Cudnn7.6.5 and Cuda 10.1 , python 3.8
```

This command shows list of physical devices available for tensorflow. You can see GPU listed here. If you have NVIDIA GPU you need to install CUDA toolkit and cuDNN as per instruction on this webpage. Without proper installation you will not see GPU in list of devices

<https://shawnhymel.com/1961/how-to-install-tensorflow-with-gpu-support-on-windows/>

```
In [5]: tf.config.experimental.list_physical_devices()
```

```
Out[5]: [PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU'),
PhysicalDevice(name='/physical_device:XLA_CPU:0', device_type='XLA_CPU'),
PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU'),
PhysicalDevice(name='/physical_device:XLA_GPU:0', device_type='XLA_GPU')]
```

```
In [6]: tf.__version__
```

```
Out[6]: '2.2.0'
```


```
In [7]: tf.test.is_built_with_cuda()
```

```
Out[7]: True
```

Load the dataset

Our dataset contains 60000 small training images that belongs to one of the below 10 classes

airplane



automobile



bird



cat



deer



dog**frog****horse****ship****truck**

In [8]: `(X_train, y_train), (X_test,y_test) = tf.keras.datasets.cifar10.load_data()`

In [9]: `X_train.shape`

Out[9]: `(50000, 32, 32, 3)`


In [10]: `y_train.shape`

Out[10]: `(50000, 1)`

Data Visualization

In [11]: `def plot_sample(index):
plt.figure(figsize = (10,1))
plt.imshow(X_train[index])`


In [12]: `plot_sample(0)`



In [13]: `plot_sample(1)`




In [14]: `plot_sample(2)`



In [15]: `classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]`

```
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

In [16]: `plot_sample(3)`



In [17]: `classes[y_train[3][0]]`

Out[17]: 'deer'

In [18]: `y_train[:3]`

Out[18]: `array([[6],
[9],
[9]], dtype=uint8)`

In [19]: `y_test.shape`

Out[19]: (10000, 1)

In [20]: `X_train.shape`

Out[20]: (50000, 32, 32, 3)

Preprocessing: Scale images

In [15]: `X_train_scaled = X_train / 255
X_test_scaled = X_test / 255`

In [16]: `y_train_categorical = keras.utils.to_categorical(
y_train, num_classes=10, dtype='float32'
)
y_test_categorical = keras.utils.to_categorical(
y_test, num_classes=10, dtype='float32'
)`

In [17]: `y_train[0:5]`

Out[17]: `array([[6],
[9],
[9],
[4],
[1]], dtype=uint8)`

In [18]: `y_train_categorical[0:5]`

Out[18]: `array([[0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.]])`

```
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
[0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
[0., 1., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```

Model building and training

In [19]:

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32,32,3)),
    keras.layers.Dense(3000, activation='relu'),
    keras.layers.Dense(1000, activation='relu'),
    keras.layers.Dense(10, activation='sigmoid')
])

model.compile(optimizer='SGD',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.fit(X_train_scaled, y_train_categorical, epochs=1)
```

```
1563/1563 [=====] - 3s 2ms/step - loss: 1.8642 -
accuracy: 0.3328
```

Out[19]:

Let's make some predictions

In []:

```
np.argmax(model.predict(X_test_scaled)[0])
```

In []:

```
y_test[0]
```

In [21]:

```
def get_model():
    model = keras.Sequential([
        keras.layers.Flatten(input_shape=(32,32,3)),
        keras.layers.Dense(3000, activation='relu'),
        keras.layers.Dense(1000, activation='relu'),
        keras.layers.Dense(10, activation='sigmoid')
    ])

    model.compile(optimizer='SGD',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model
```

Measure training time on a CPU

In [22]:

```
%timeit -n1 -r1
with tf.device('/CPU:0'):
    cpu_model = get_model()
    cpu_model.fit(X_train_scaled, y_train_categorical, epochs=1)
```

```
1563/1563 [=====] - 44s 28ms/step - loss: 1.8660
- accuracy: 0.3301
44.5 s ± 0 ns per loop (mean ± std. dev. of 1 run, 1 loop each)
```

Lets measure training time on a GPU (I've NVIDIA

Titan RTX)

```
In [23]: %%timeit -n1 -r1
          with tf.device('/GPU:0'):
              cpu_model = get_model()
              cpu_model.fit(X_train_scaled, y_train_categorical, epochs=1)
```

1563/1563 [=====] - 3s 2ms/step - loss: 1.8581 - accuracy: 0.3354
 3.6 s ± 0 ns per loop (mean ± std. dev. of 1 run, 1 loop each)

Lets run same test for 10 epocs

```
In [25]: %%timeit -n1 -r1
          with tf.device('/CPU:0'):
              cpu_model = get_model()
              cpu_model.fit(X_train_scaled, y_train_categorical, epochs=10)
```

Epoch 1/10
 1563/1563 [=====] - 47s 30ms/step - loss: 1.8614 - accuracy: 0.3328
 Epoch 2/10
 1563/1563 [=====] - 48s 30ms/step - loss: 1.6600 - accuracy: 0.4130
 Epoch 3/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.5713 - accuracy: 0.4486
 Epoch 4/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.5106 - accuracy: 0.4663
 Epoch 5/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.4596 - accuracy: 0.4837
 Epoch 6/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.4157 - accuracy: 0.5000
 Epoch 7/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.3778 - accuracy: 0.5155
 Epoch 8/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.3429 - accuracy: 0.5262
 Epoch 9/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.3088 - accuracy: 0.5392
 Epoch 10/10
 1563/1563 [=====] - 44s 28ms/step - loss: 1.2786 - accuracy: 0.5497
 7min 26s ± 0 ns per loop (mean ± std. dev. of 1 run, 1 loop each)

```
In [24]: %%timeit -n1 -r1
          with tf.device('/GPU:0'):
              cpu_model = get_model()
              cpu_model.fit(X_train_scaled, y_train_categorical, epochs=10)
```

Epoch 1/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.8624 - accuracy: 0.3322
 Epoch 2/10

```
1563/1563 [=====] - 3s 2ms/step - loss: 1.6579 -  
accuracy: 0.4146  
Epoch 3/10  
1563/1563 [=====] - 3s 2ms/step - loss: 1.5686 -  
accuracy: 0.4443  
Epoch 4/10  
1563/1563 [=====] - 3s 2ms/step - loss: 1.5068 -  
accuracy: 0.4681  
Epoch 5/10  
1563/1563 [=====] - 3s 2ms/step - loss: 1.4572 -  
accuracy: 0.4863  
Epoch 6/10  
1563/1563 [=====] - 3s 2ms/step - loss: 1.4139 -  
accuracy: 0.5009  
Epoch 7/10  
1563/1563 [=====] - 3s 2ms/step - loss: 1.3766 -  
accuracy: 0.5147
```